

DL Pac 3B

```
In [2]: import tensorflow as tf
import matplotlib.pyplot as plt
from tensorflow import keras
import numpy as np
(x_train, y_train), (x_test, y_test) = keras.datasets.fashion_mnist.load_data()
```

WARNING:tensorflow:From C:\Users\rushi\anaconda3\lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz>)

29515/29515 [=====] - 0s 0s/step

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz>)

26421880/26421880 [=====] - 8s 0us/step

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz>)


5148/5148 [=====] - 0s 0s/step

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz>)

4422102/4422102 [=====] - 1s 0us/step


```
In [3]: plt.imshow(x_train[1])
```

```
Out[3]: <matplotlib.image.AxesImage at 0x27b68d9e820>
```



```
In [4]: plt.imshow(x_train[0])
```

```
Out[4]: <matplotlib.image.AxesImage at 0x27b68e1f850>
```



```
In [6]: # Next, we will preprocess the data by scaling the pixel values to be between 0 and 1, and then
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0
x_train = x_train.reshape(-1, 28, 28, 1)
x_test = x_test.reshape(-1, 28, 28, 1)
```

```
In [8]: x_train.shape
```

```
Out[8]: (60000, 28, 28, 1)
```

```
In [9]: x_test.shape
```

```
Out[9]: (10000, 28, 28, 1)
```

```
In [10]: y_train.shape
```

```
Out[10]: (60000,)
```

```
In [11]: y_test.shape
```

```
Out[11]: (10000,)
```

```
In [14]: model = keras.Sequential([
    keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    keras.layers.MaxPooling2D((2, 2)),
    # It shows a 13 * 13 size image with 32 channels or filters or depth.
    keras.layers.Dropout(0.25),
    # Reduce Overfitting of Training sample drop out 25% Neuron
    keras.layers.Conv2D(64, (3, 3), activation='relu'),
    # 64 * 3 * 3 = 576 + 1 = 577 * 32 + 32(bias) = 18496
    keras.layers.MaxPooling2D((2, 2)),
    # It shows a 5 * 5 size image with 64 channels or filters or depth.
    keras.layers.Dropout(0.25),
    keras.layers.Conv2D(128, (3, 3), activation='relu'),
    # We need to flatten the 3x3x128 feature map to a vector of size 1152
    keras.layers.Flatten(),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dropout(0.25),
    keras.layers.Dense(10, activation='softmax')
])

model.summary()
```

WARNING:tensorflow:From C:\Users\rushi\anaconda3\lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From C:\Users\rushi\anaconda3\lib\site-packages\keras\src\layers\pooling\max_pooling2d.py:161: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
dropout (Dropout)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
dropout_1 (Dropout)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 128)	73856
flatten (Flatten)	(None, 1152)	0
dense (Dense)	(None, 128)	147584
dropout_2 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 10)	1290
=====		
Total params: 241546 (943.54 KB)		
Trainable params: 241546 (943.54 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
In [ ]: # Compile and Train the Model
# After defining the model, we will compile it and train it on the training data.
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
history = model.fit(x_train, y_train, epochs=10, validation_data=(x_test, y_test))
```

```
In [18]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
print('Test accuracy:', test_acc)
```

```
313/313 - 1s - loss: 0.2541 - accuracy: 0.9101 - 1s/epoch - 5ms/step
Test accuracy: 0.910099983215332
```